

Unsupervised Representation learning by predicting image rotations

①

Summary

1. Train Convnet to predict the 2d rotation that is applied to the image it gets as input.

2. Demonstrate dramatic perf improvement compared to prior UL methods.

3. 2 days on a single Titan X GPU.

Predicting Rotations

1. 0, 90, 180, 270 degrees.

2. The ConvNet has to predict the transformation without having access to the original untransformed image.

3. Intuitive argument: it is not possible to perform this task without knowing how to:

- recognize classes of the objects in the image
- ~~the~~ detect the semantic parts of the object

4. More concretely, the model needs to:

- localize salient objs in the imgs
- recognize their orientation and type
- relate the obj's orientation with the dominant orientation that each type of objects tend to be depicted within the available images.

5. Argument w/ prior methods: image rotations do not create low-level visual artifact that allow the network to cheat in the task